

## Scope (from the diagram)

- **Your START:** when the **Market Analysis Agent** (fed by Search/Data-Collection) emits a **company\_schema v1.0 JSON record**.
- **Your END:** a persisted **scoring payload** (score + breakdown + reason text + risk/feasibility flags), exposed via **REST** to:
  - **Opportunity Validation Agent** (gate/triage),
  - **Report Generation Agent** (the “why” panel),
  - **Results DB / Memory** (audit, cohort stats).

Not your scope: agent orchestration, crawling/scraping, NLP extraction, knowledge graph, UI.

---

## Phase 0 — Contracts & Scaffolding (your entry handshake)

**Goal:** Make the inputs/outputs unambiguous so upstream and downstream can integrate without you.

**Entrypoint:** company\_schema v1.0 (from Market Analysis Agent).

```
{
  "company_id": "lawfirm-123",
  "domain": "examplelaw.com",
  "digital": { "pagespeed": 78, "crm_flag": true, "ecom_flag": false },
  "ops": { "employees": 35, "locations": 2, "services_count": 6 },
  "info_flow": { "daily_docs_est": 240 },
  "market": { "competitor_density": 14, "industry_growth_pct": 6.1, "rivalry_index": 0.42 },
  "budget": { "revenue_est_usd": 4200000 },
  "meta": { "scrape_ts": "2025-08-07T17:15:00Z", "source_confidence": 0.83 }
}
```

**Exit:** validator + example fixtures everyone can test against.

#### What you build

- **company\_schema v1.0** (strict types/ranges/null policy) and a JSON-Schema/Pydantic model.
- **weights.yaml v0** with the five weights + comments on meaning and change control.
- **Fixture set** (JSONL) with realistic and edge-case companies.
- **CI contract tests** that fail on schema drift.

#### Done when

- Upstream can run `validate --in companies.jsonl` and see clear errors/success.
  - You've frozen `company_schema@1.0.0` (minor bumps allowed, with tests).
- 

## Phase 1 — Normalization & Cohort Stats

**Goal:** Turn raw fields into **unitless, comparable** features deterministically.

**Entrypoint:** validated company records.

**Exit:** normalized feature vector + a **NormContext** (mean/ $\sigma$  snapshots).

#### What you build

- zscore, log10p, flag coercion, clipping, null policy (documented).
- **Cohort stats capture** (mean/std per field) for each batch; ID the stats used (e.g., `norm_stats_id`).
- Strict handling of missing/low-confidence inputs (warnings, never silent guesses).

#### Done when

- Re-running the same batch with the same NormContext yields identical features.

---

## Phase 2 — Sub-scores (D/O/I/M/B)

**Goal:** Compute the five sub-scores exactly as spec + explanations.

**Entrypoint:** normalized features + flags.

**Exit:** {value  $\in$  [0,1], inputs\_used, warnings} for each of D/O/I/M/B.

**What you build**

- **D (Digital Maturity):**  $0.4 * \text{PageSpeed\_scaled} + 0.3 * \text{CRM} + 0.3 * \text{Ecom}$ .
- **O (Ops Complexity):**  $z(\text{employees}) + z(\text{locations}) + z(\text{services})$ .
- **I (Info Flow):**  $\log_{10}(\text{daily\_docs} + 1) / 4$ .
- **M (Market Pressure):**  $z(\text{comp\_density}) + z(\text{industry\_growth}) - z(\text{rivalry})$ .
- **B (Budget):**  $\log_{10}(\text{revenue}) / 7$ .
- Each returns value + what drove it (inputs/fallbacks).

**Done when**

- Golden tests (hand-calc'd) match within 1e-6 and every output has explain bits.

---

## Phase 3 — Opportunity Score + Explainability

**Goal:** Produce the 0–100 score and a human-readable **why**.

**Entrypoint:** 5 sub-scores + weights.

**Exit:** scoring payload for a single record.

**What you build**

- **Final:** Score =  $100 \times (0.25D + 0.20O + 0.20I + 0.20M + 0.15B)$  with consistent rounding.
- **Contributions:** points = weight\*value\*100, sorted; top drivers.
- **Reason text (templated, deterministic):** e.g., “High digital maturity (fast site, CRM) and strong document volume; budget moderate.”
- **Warnings passthrough** (e.g., missing revenue treated as 0).

#### Done when

- Random samples read correctly to a human and match the numbers shown.

---

## Phase 4 — Risk & Feasibility Gates (lightweight validators)

**Goal:** Provide triage signals to the **Opportunity Validation Agent**.

**Entrypoint:** scoring payload + input confidences.

**Exit:** risk and feasibility blocks attached to the payload.

#### What you build

- **Risk:** data confidence score, missing-field penalties, scrape volatility.
- **Feasibility:** boolean gates for “deployable now” (docs present, CRM/e-com/booking present, budget above floor).
- Clear reasons when a company is flagged.

#### Done when

- Validation Agent can filter: {score >= X AND feasibility.ok == true AND risk.low}.
-

# Phase 5 — Batch Runner & Persistence

**Goal:** Score cohorts reproducibly and persist audit-ready outputs.

**Entrypoint:** JSONL of companies.

**Exit:** JSONL of scores + DB rows + a stored NormContext.

## What you build

- Batch CLI: `score --in companies.jsonl --out scores.jsonl`.
- **NormContext persistence** (stats snapshot + ID).
- **Results DB writes** (score, breakdown, reason, risk, feasibility, norm\_stats\_id).

## Done when

- Same input + same NormContext  $\Rightarrow$  identical output file checksum and DB rows.
- 

# Phase 6 — Scorer Service (REST) + Integration Points

**Goal:** Let downstream pieces call you synchronously/asynchronously.

**Entrypoint:** HTTP requests from agents/UI/backends.

**Exit:** JSON responses + webhooks (optional) + metrics.

## What you build

- **FastAPI endpoints:**
  - POST `/score` (single record), POST `/score/batch` ( $\leq 1k$ ), GET `/healthz`, GET `/stats` (last NormContext).
- **Schema validation** with precise error messages.

- **OpenAPI** with copy-paste examples for Report Gen & Validation Agents.

**Done when**

- Contract tests pass; p99 <150ms/record locally; clean error semantics.
- 

## Phase 7 — Back-testing & Calibration (freeze the math)

**Goal:** Show the score tracks outcomes; tune weights if helpful.

**Entrypoint:** 200-SMB labeled cohort + your scores.

**Exit:** calibration report + frozen weights.yaml v1.0.

**What you build**

- Metrics: Spearman/Kendall rank corr, AUC (if binary), Precision@K/Lift.
- One-pager with plots and a go/no-go on weight tweaks.
- Tag the release (weights@1.0, scorer@1.0.0).

**Done when**

- Evidence is clear, weights are frozen, and results are reproducible.
- 

## Phase 8 — Sensitivity, Robustness & Drift

**Goal:** Prove stability and catch input drift early.

**Entrypoint:** baseline run.

**Exit:** stability report + alerts.

### What you build

- **$\pm 10\%$  weight sweeps** (sum preserved)  $\rightarrow$  Kendall  $\tau$  vs. baseline; top-K overlap.
- **Input drift checks** (mean/ $\sigma$  deltas, null-rate alarms).
- Thresholds that page you only when it actually matters.

### Done when

- Stability documented; drift alarms tested; thresholds agreed with the lead.
- 

## Phase 9 — Observability & QA (operational hardening)

**Goal:** Make it safe to run without hand-holding.

**Entrypoint:** service + batch jobs.

**Exit:** actionable logs/metrics and failure modes that are obvious.

### What you build

- Structured logs (latency, warning rates, nulls), histograms per field.
- Fuzz/mutation tests on formulas and schema.
- SLOs + light runbook (deploy, rollback, change control for weights/schema).

### Done when

- A bad payload fails fast with a helpful message; on-call can triage from logs alone.
-

# Phase 10 — Downstream Enablement (your exit handshake)

**Goal:** Make your output drop-in for the **Opportunity Validation** and **Report Generation** Agents.

**Entrypoint:** your REST and DB rows.

**Exit:** tiny SDKs/docs; no extra logic needed elsewhere.

## What you build

- Mini-SDK helpers (score\_record, score\_batch) + curl examples.
- Output examples the Report Gen Agent can render as-is (score, breakdown, contributions, reason, risk/feasibility).
- “Hello-world” doc to score a JSONL in <15 minutes.

## Done when

- Both agents call you in their pipelines without requiring changes to your code.
- 

## Quick ownership map (one line each)

- **In from:** Market Analysis Agent → company\_schema v1.0.
- **You:** Normalize → Sub-scores → Score (0–100) → Explain → Risk/Feasibility → Persist → Serve.
- **Out to:** Opportunity Validation Agent, Report Generation Agent, Results DB/Memory.

Want me to spit out the OpenAPI for POST /score and a company\_schema v1.0 draft next so the agent team can wire against it?